

# Sistemas Distribuídos

José Orlando Pereira


Departamento de Informática  
Universidade do Minho





# Clock synchronization

- Hardware clocks are not perfect and drift over time
- Clock skew can be a problem with:
  - shared files, “make”, ...
  - database systems (e.g. Google Spanner)

Ideally:




# Clock synchronization



- The clock must be adjusted in small increments, over a longer period of time by making it faster or slower

# Clock synchronization


- In practice, there are unpredictable
  - transmission delays
  - processing delays



- The best we can do is adjust by  $(T_2 - T_1) - \text{(estimated message delay)}$

# Clock synchronization


- What is the message delay?



- Network Time Protocol (NTP):
  - assume delays are the same =  $((T_2-T_1)+(T_4-T_3))/2$
  - repeat several times and pick the smallest delay

# Clock synchronization

- Reference Base Synchronization (RBS):
  - assume true broadcast medium (aprox. zero delay)




# Mutual exclusion

- Solution in a distributed system?
- Recall the definition of mutual exclusion:
  - No two threads in the critical section
  - No deadlock / no starvation
- We consider:
  - Number of message hops to entering critical section
  - Load balancing


# Mutual exclusion

- Centralized queue kept by a coordinator:
  - 1 round-trip to enter / asymmetric load




# Mutual exclusion

- Exchange a token in a ring:
  - $n/2$  hops to entering / symmetric load




# Mutual exclusion

- A distributed algorithm is hard to achieve:
  - As concurrent lock requests are received by different destinations in different orders, safety is not ensured




# Mutual exclusion

- Taking advantage of synchronized clocks:
  - Assume  $\delta > (\text{transmission delay} + \text{skew})$
  - Consider only messages up to  $t-\delta$ , order by timestamp




# Timeouts

- Used to assess status of remote processes
- Tight timeouts are dangerous:
  - E.g., proportional to mean delay
  - Means low coverage
- Large timeouts are not useful:
  - E.g., proportional to high percentile
  - Taking advantage of time causes a very large performance penalty



# Using real time

- Solutions that do not use time are more robust:
  - In wide area networks
  - With performance perturbations
- Solutions that do not use time might have better performance:
  - Run time proportional to mean delay
  - Even if more message exchanges are necessary




# Asynchronous system model

- Assume no global time reference
- Assume no bounds on:
  - clock drift
  - processing time
  - message passing time
- Can we still solve important problems?

# Time and causality

- What is special about time that makes it useful for distributed algorithms?




# Time and causality

- $\text{Clock}(i)$  the time at which i happened
- If  $i$  precedes  $j$  then  $\text{Clock}(i) < \text{Clock}(j)$
- For some event j:
  - When we are sure that there is no unknown i such that  $\text{Clock}(i) < \text{Clock}(j)$
  - Then there is no i such that  $i$  precedes  $j$
- **Can we build a logical clock with the same property?**

# Lamport's logical clocks

- Local events: increment counter
- Send events: increment and then tag with counter
- Receive events: update local counter to maximum and then increment




# Mutual exclusion

- Algorithm sketch:
  - Start by assuming that processes are continually exchanging messages over FIFO channels...
  - $r_i[j]$  latest timestamp from  $j$  at  $i$
  - Consider requests with  $t \leq \min(r_i[j], \text{for all } j)$ 
    - (akin to  $t-\delta!$ )
  - Order by timestamp, break ties by process id
- (The complete version is the Ricart-Agrawala distributed mutex algorithm)

# Mutual exclusion

- 1 hop to enter / symmetric load



# Conclusion

- The same approach used for the waiting queue in the mutex can be used for other deterministic applications
  - Replicated State Machine (RSM)
- Logical time is widely applicable in distributed systems to solve many problems