

Sistemas Distribuídos

José Orlando Pereira

Departamento de Informática
Universidade do Minho




Fault tolerance


- A distributed system is composed by autonomous computing elements and can tolerate faults to avoid failure
- Fault model describes types of faults:
 - Omissive: process crash, lost message, ...
 - Assertive (a.k.a. Byzantine): corrupted messages, ...
- Fault model describes the number of faults:
 - Example: Number of processes that can crash

Transactional commit


- Coordinate multiple irreversible actions across a distributed system:




2-phase commit (2PC)




2PC: Participant failure




2PC: Coordinator failure




2PC in systems



Transactional RPC



Transactional RPC




Summary

- Distributed transactions with 2-phase commit (2PC) support agreement in systems with faults
 - Limited to crash-recovery of the coordinator
- Is widely used in enterprise middleware for application integration

Replication

- Keep multiple copies of the same data or service
 - Distribute the load for scalability
 - Tolerate server faults




Replication

- Naive solution: write then propagate
 - state may diverge
 - clients observe paradoxes when reading
 - not fault-tolerant
- 2PC: Correct, but progress only with all up (tolerates reboots, not crashes)




Replication

- Assume that:
 - operations are reads and writes
 - we keep a timestamp with each item
- It might be possible to read and write from fewer processes...



Quorum

- Assume 2-phase protocol for writing (phase 1 == read)
- Quorum rules for replicated data:
 - $N_R + N_W > N \rightarrow$ readers get the latest value
 - $N_W > N/2 \rightarrow$ concurrent writers conflict



Quorum

- Additional rules for fault-tolerance when assuming at most f faults:
 - $N_R + f \leq N \rightarrow$ readers never block
 - $N_W + f \leq N \rightarrow$ writers never block
- Can be configured to ensure both or either of them
- Typical solution is having a majority:
 - $N_R = N_W = f + 1$
 - $N = 2f + 1$
- Examples: $N=3$ for $f=1$, $N=5$ for $f=2$, ...

Summary

- Flexible and efficient solution for data replication
 - Example: Amazon Aurora DB
- Wasteful when there are multiple concurrent writers:
 - At most one of multiple write operation can be accepted
 - But it can happen that none is accepted if each operation is applied in less than N_w servers

Leader election (Bully algorithm)

- A process that wants to be the leader (suspects that there is no leader), broadcasts its address to all others
- If address > local is received, acknowledge the new leader
- If address < local is received, broadcast local address (i.e. bully other candidates into submission)
- Transiently, there can be multiple or no leaders...


Consensus

- Distributed agreement problem, in which processes propose alternatives and decide as follows:
 - All correct participants get a decision
 - All decisions are the same
 - The decision is one of the proposed alternatives
- How to solve this problem?
- 2PC + Quorum + Bully $\sim =$ Paxos consensus!

Paxos

- Processes in 3 main roles:
 - Proposers (with a Leader) → coordinate the voting
 - Acceptors → vote for the decision
 - Learners → are informed of the decision
- Usually, each participant plays all the 3 roles

Paxos



Paxos

- Phase 1 is leader election + quorum read
 - Propose (1a):
 - A proposer tries to become the current leader
 - Uses (time,id) to bully others
 - Promise (1b):
 - The acceptor promises to forget previous leaders
- A Leader is established with a majority of promises and uses the value read or, if none, its own opinion as the candidate value

Paxos

- Phase 2 is quorum write:
 - Accept (2a):
 - The leader imposes its decision
 - Accepted (2b):
 - Accepted if the Leader was not deposed
 - The acceptor confirms the decision to Leader and Learners
- Upon receiving a majority of accepted messages, the decision is done!

Paxos

- Can always recover from a minority of failed processes if there is a trusted leader
- What's the worst case scenario?
 - No trusted Leader (unstable system)
 - More than a minority of faults (catastrophic fault)
- In both cases, the algorithm stops but never produces inconsistent decisions

Applications of consensus

- Replicated State Machine: Replicas execute the same sequence of commands
 - Clients issue commands to Proposers
 - Learners execute each request decided as the next and send replies to Clients
- Group communication: Agreement on operational processes that receive each message
 - The decision is on group membership and messages seen by each participant

Summary

- Consensus protocols at the heart of fault-tolerant distributed systems:
 - Container orchestration in Kubernetes
 - Transactions in Google Spanner database
- Consensus algorithms are available for even more demanding scenarios (Byzantine faults), where participants can misbehave, lie, ...
 - Applicable to blockchains